

Module 2

Designing Cloud Applications for
Resiliency

Module Overview

- Application Design Practices for Highly Available Applications
- Application Analytics
- Building High Performance Applications by Using ASP.NET
- Common Cloud Application Patterns
- Caching Application Data

Lesson 1: Application Design Practices for Highly Available Applications

- Partitioning Workloads
- Load Balancing
- Transient Fault Handling
- Queues

Partitioning Workloads

- When designing web applications, split your business processes into Partitioning Workloads
- Partitioning Workloads:
 - Can be handled in modular websites, cloud services, or virtual machines
 - Provides the ability to scale the different components of your application in isolation

Partitioning Workloads (cont.)

- Example Photo Sharing Application

Virtual Machine

Web Front-End

SignalR Hub

Image Processor

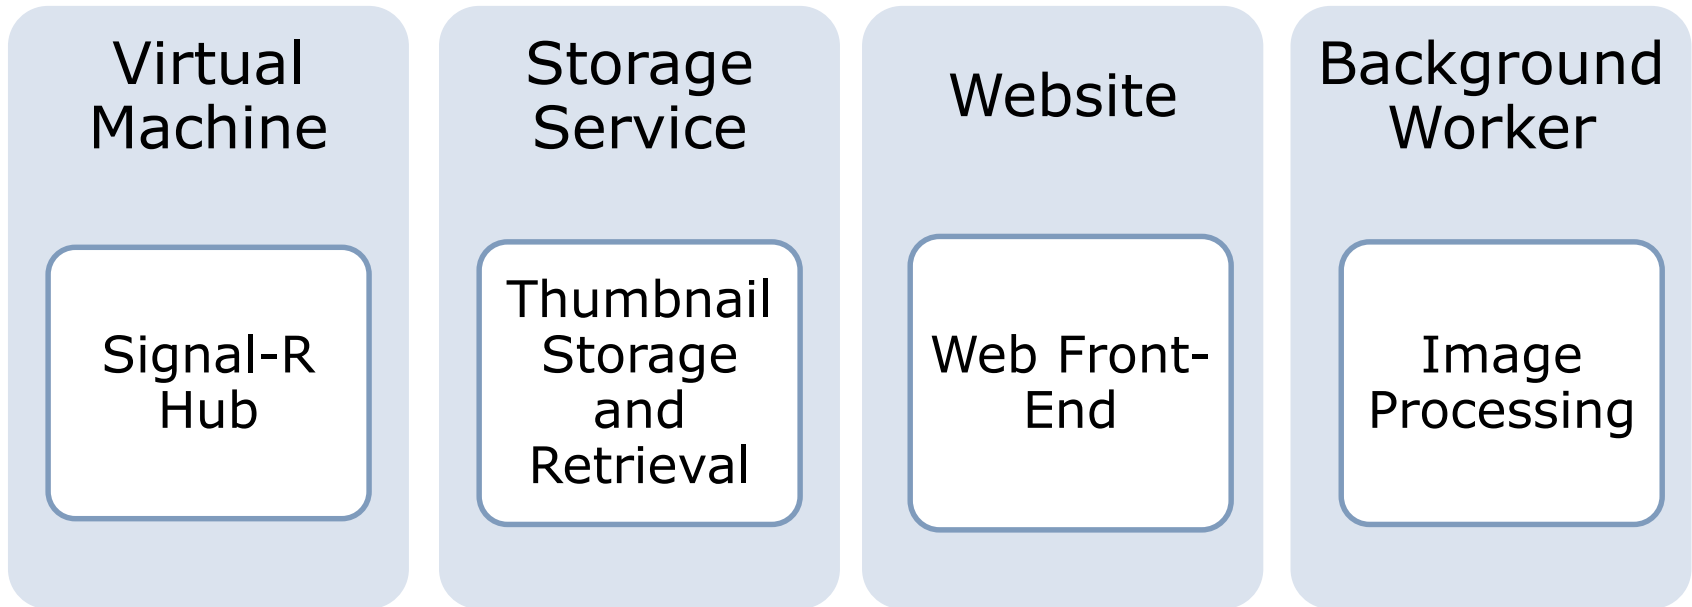
Thumbnail Storage

Partitioning Workloads (cont.)

- Example Photo Sharing Application
 - The Web Front-End shows thumbnails of images users have uploaded today
 - The application code takes an image that a user uploads, processes it into a thumbnail and then stores the thumbnail.
 - The SignalR hub notifies the client web browser that the image is finished processing

Partitioning Workloads (cont.)

- Example Photo Sharing Application



- Each component of the application can be scaled in isolation to meet its specific demand

Load Balancing

- Provide the same service from multiple instances and use a load balancer to distribute requests across all of the instances
- Considerations for selecting a load balancing strategy:
 - Hardware or software load balancers
 - Load balancing algorithms (round robin)
 - Load balancer stickiness

Load Balancing and Geographic Resiliency

- Load balancing can be combined with geographic redundancy to help achieve high availability
 - You can use a load balancer to direct client requests to their closest data center
 - Traffic Manager is often used for this task
 - A load balancer can be used to implement a failover scenario
 - When a data center or compute instance is down, clients can be directed to the next desirable instance that is available
- Designing a load balancing strategy along with distributing your application across data centers is key to high availability across the globe

Transient Fault Handling

- Transient faults can occur because of a temporary condition
 - Service is unavailable
 - Network connectivity issue
 - Service is under heavy load
- Retrying your request can resolve temporary issues that normally would crash an application
- You can retry using different strategies
 - Retry after a fixed time period
 - Exponentially wait longer to retry
 - Retry in timed increments

Transient Fault Handling (cont.)

- The Transient Fault Handling application block is a part of the Enterprise Library
- The Enterprise library contains a lot of code that is necessary to implement the pattern and the retry strategies
- Retry strategies are prebuilt for common scenarios including accessing Azure services
- You can build custom strategies and extend the library for the unique needs of your application

Queues

- A modular web application can behave like a monolithic application if each component relies on a direct two-way communication with a persistent connection
- Persistent queue messages allow your application to handle requests if one of your application components fail or is temporary unavailable
- An external queue allows your application to audit requests or measure your load without adding any overhead to the code of your primary application

Queues (cont.)

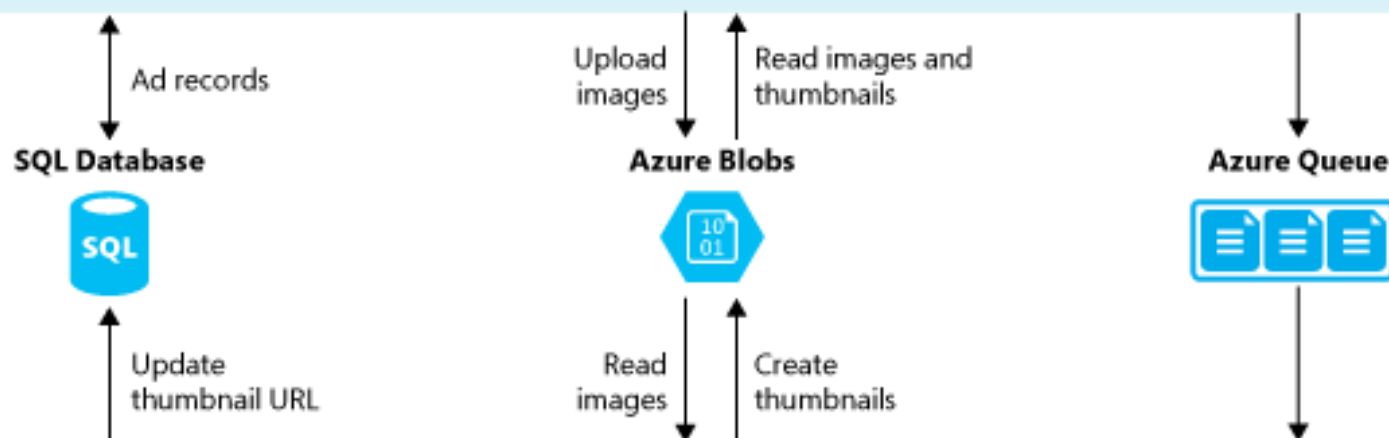
- Queues can be used to communicate in a disconnected manner between the different components of your application
 - If an instance of your application module fails, another instance can process the queue messages
 - Messages that fail to be processed can either be reinserted into the queue or not marked as complete by the client module
 - Messages can be measured or audited in isolation from your application
- Queues become very important when dealing with background processing (Worker roles and WebJobs)

Queues (cont.)



Web Role

- Users place ads by entering ad info and uploading an image. The app writes a database record, stores an image blob, and writes a queue message to schedule image processing to create a thumbnail.
- Users view ads by selecting a category and optionally entering a search string. The app reads the database for ad info and for image and thumbnail blob URLs.



Worker Role

- Polls images queue for images to be processed.
- When a new queue message is received:
 - Converts the image to a thumbnail, saves blob
 - Adds thumbnail URL to ad record in database
 - Deletes the queue message

Lesson 2: Application Analytics

- Application Insights
- Demonstration: Monitoring a Web Application

Application Insights

- Application Insights is a analytics and monitoring service available for your applications
 - View exception stack traces
 - Monitor CPU and resource usage
 - Periodically test URLs from worldwide data centers
 - Monitor usage of your application and most popular requests
- It can be used with .NET or Java
- Applications do not specifically need to be hosted in Azure

Application Insights (cont.)



- The application insights blade provides a rich dashboard of data (tiles) and supports “drilling down” into more specific metrics.

Demonstration: Monitoring a Web Application

In this demonstration, you will learn how to:

- Create a new ASP.NET Web Application project with Application Insights referenced
- Use the Application Insights service and dashboard to monitor the metadata about your application's requests

Lesson 3: Building High Performance Applications by Using ASP.NET

- Web Hosting Performance Improvements
- Asynchronous HTTP Modules and Handlers
- The Async Keyword
- State Management

Web Hosting Performance Improvements

- ASP.NET 4.5 introduced many improvements to web hosting:
 - Asynchronous requests and responses
 - Support for Task Parallel Library and await and async keywords
 - Sharing common assemblies across applications
 - Multi-core startup (JIT)
 - Windows prefetcher for web applications

Asynchronous HTTP Modules and Handlers

- Asynchronous methods in ASP.NET allow your code that is waiting for an IO-bound operation (Database, Service, Disk) to return threads to the ThreadPool
 - Typically these requests would block the thread from servicing other requests
 - By releasing threads while waiting on external factors, the amount of requests that can be processed simultaneously increases exponentially
- Helps greatly with applications where the load spikes and is not generally consistent

The Async Keyword

- The **async** and **await** keywords are available to create easy to read and write asynchronous methods in C#
- The keywords can be used with ASP.NET MVC to create asynchronous actions:

```
public async Task<ActionResult> ItemsAsync() {  
    var context = new DatabaseContext();  
    var model = await context.GetModelItemsAsync();  
    return View("Items", model);  
}
```

State Management

- When distributing your application across multiple instances, session state needs to be shared across the instances
 - What happens when a user starts on Server1 but is now on Server2 when they click a link?
- Session state can be moved from in-memory to a dedicated session server:
 - Microsoft SQL Server
 - ASP.NET State Server
- Session state can also be partitioned and distributed among multiple session stores

Lesson 4: Common Cloud Application Patterns

- Cloud Application Patterns
- The Retry Pattern
- The Valet Key Pattern


Cloud Application Patterns

- MSDN provides a collection of cloud design patterns
- Patterns:
 - Retry
 - Valet Key

The Retry Pattern

- The Retry pattern is designed to handle temporary failures
- Failures are assumed to be transient until they exceed the retry policy
- The Transient Fault Handling Block is an example of a library that is designed to implement the Retry pattern (and more)
- Entity Framework provides a built-in retry policy implementation
 - Implemented in version 6.0

The Retry Pattern (cont.)

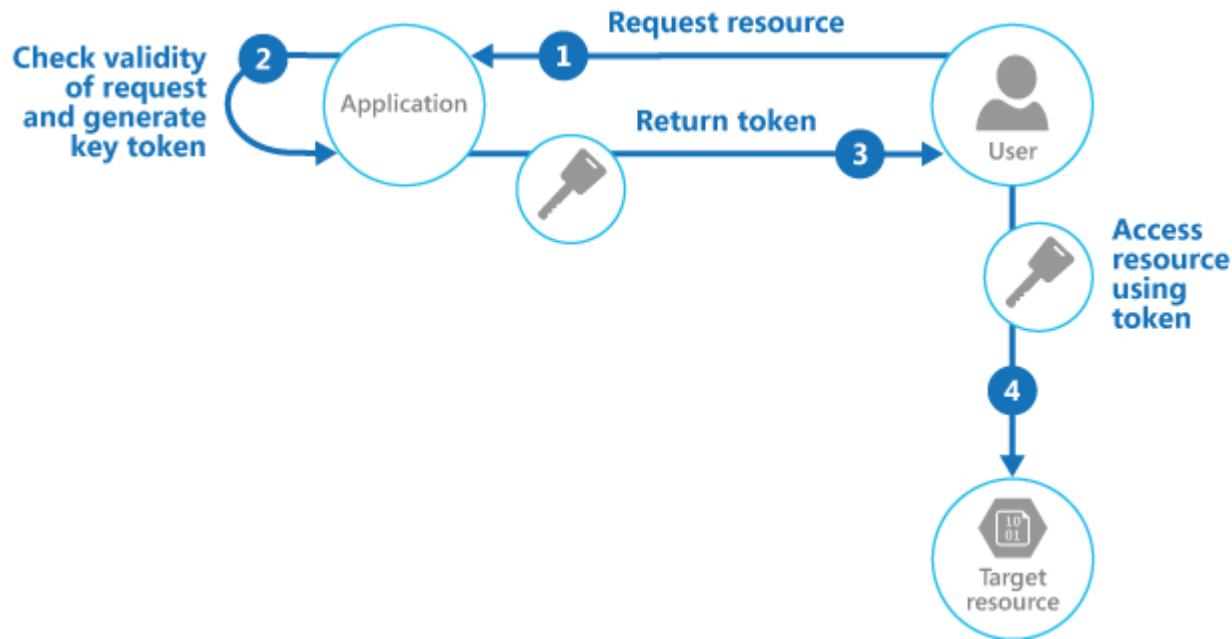


- The application sends a request to a hosted service
- The hosted service responds with a HTTP 500 (Internal Server Error) code
- The application retries until it exceeds the retry count for its policy or is met with an acceptable status code such as HTTP 200 (OK)

The Valet Key Pattern

- If your application access a resource on behalf of your clients, your servers take on additional load
- You do not want to make your resources publically available so you are typically forced to have your application validate a client
- The Valet Key pattern dictates that your application simply validates the client and then returns a token to the client. The client can then retrieves the resource directly using its own hardware and bandwidth

The Valet Key Pattern (cont.)



- The client requests a resource from your application
- Your application validates the client and then returns an access token
- The client then directly accesses the resource by using the provided token

Lesson 5: Caching Application Data

- Redis Cache

Redis Cache

- Based on the open-source Redis platform
 - Multiple tiers are available that offer different numbers of nodes
 - Supports transactions
 - Supports message aggregation using a publish subscribe model
 - Considered a key-value store where the keys can be simple or complex values
 - Massive Redis ecosystem already exists with many different clients

Module Review and Takeaways

- Review Question(s)
- References:
 - <https://docs.microsoft.com/en-us/azure/app-service/choose-web-site-cloud-service-vm>
 - <https://docs.microsoft.com/en-us/azure/cloud-services/cloud-services-dotnet-get-started>
 - <https://github.com/Azure/azure-webjobs-sdk/wiki>
 - <https://docs.microsoft.com/en-us/azure/application-insights/>
 - <https://docs.microsoft.com/en-us/azure/application-insights/app-insights-snapshot-debugger>
 - <https://docs.microsoft.com/en-us/azure/redis-cache/cache-dotnet-how-to-use-azure-redis-cache>